

## Apprentissage Artificiel

## Logistic Regression

Ewa Kijak

ESIR/Univ. Rennes

## Outline

Introduction

Principle

Learning by optimization

Going further

## Sommaire

Introduction

Principle

Learning by optimization

Going further

## Logistic regression is classification

- ▶ Logistic regression  $\neq$  Linear regression
- ▶ Logistic regression is a Generalized Linear Model (GLM)

## Univariate linear regression

Example : Housing Prices

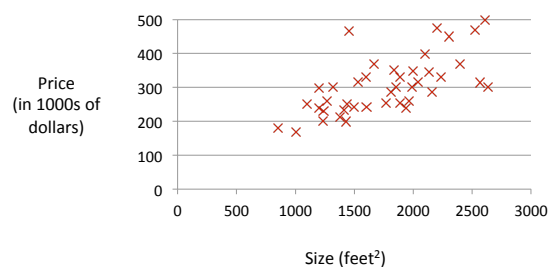
- ▶  $m$  = Number of training examples
- ▶  $x^{(i)}$  = "input" variable / features of the  $i$ -th example
- ▶  $y^{(i)}$  = "output" variable / "target" variable of the  $i$ -th example

| Size in feet <sup>2</sup> (x) | Price (\$) in 1000's (y) |
|-------------------------------|--------------------------|
| 2104                          | 460                      |
| 1416                          | 232                      |
| 1534                          | 315                      |
| 852                           | 178                      |
| ...                           | ...                      |

Introductory examples borrowed from Andrew Ng

## Univariate linear regression

- ▶ **Supervised learning** : Give the "right answer" for each example in the data.
- ▶ **Regression problem** : Predict real-valued output
- ▶  $\neq$  **Classification problem** : Discrete-valued output



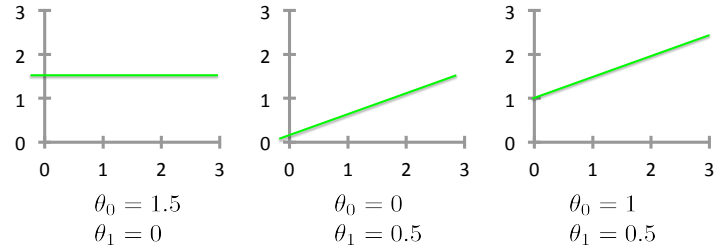
## Univariate linear regression

- Univariate linear regression = Linear regression with one variable ( $x$ ).
- Hypothesis :  $h_{\theta}(x) = \theta_0 + \theta_1 x$
- Parameters :  $\theta_0, \theta_1$

How to choose  $\theta_i$  ?

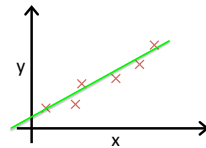
## Univariate linear regression

$$h_{\theta}(x) = \theta_0 + \theta_1 x$$



## Univariate linear regression

Idea : Choose  $\theta_0, \theta_1$  so that  $h_{\theta}(x)$  is close to  $y$  for the training examples  $(x^{(i)}, y^{(i)})$



- Cost function :

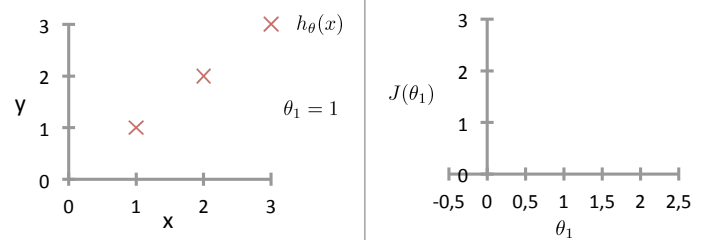
$$J(\theta_0, \theta_1) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

- Goal :

minimize  $J(\theta_0, \theta_1)$   
 $\theta_0, \theta_1$

## Univariate linear regression

Simplified case (example) :  $h_{\theta}(x) = \theta_1 x$



## Multivariate linear regression

- Multivariate linear regression = Linear regression with multiple features ( $x_j$ ).
- $d$  = number of features
- $x^{(i)}$  = "input" variable / features of the  $i$ -th example
- $x_j^{(i)}$  = value of feature  $j$  in the  $i$ -th example

Example : Housing Prices

| Size (feet <sup>2</sup> ) | Number of bedrooms | Number of floors | Age of home (years) | Price (\$1000) |
|---------------------------|--------------------|------------------|---------------------|----------------|
| 2104                      | 5                  | 1                | 45                  | 460            |
| 1416                      | 3                  | 2                | 40                  | 232            |
| 1534                      | 3                  | 2                | 30                  | 315            |
| 852                       | 2                  | 1                | 36                  | 178            |
| ...                       | ...                | ...              | ...                 | ...            |

## Multivariate linear regression

- Hypothesis :  $h_{\theta}(\mathbf{x}) = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \dots + \theta_d x_d$
- For convenience of notation, we define  $x_0 = 1$ .

Then :

- $\mathbf{x} = [x_0, x_1, x_2, \dots, x_d]^T \in \mathbb{R}^{d+1}$
- $\boldsymbol{\theta} = [\theta_0, \theta_1, \theta_2, \dots, \theta_d]^T \in \mathbb{R}^{d+1}$

$$h_{\theta}(\mathbf{x}) = \boldsymbol{\theta}^T \mathbf{x}$$

## Sommaire

Introduction

Principe

Learning by optimization

Going further

## Logistic regression

► Linear classification model

► Let  $\mathbf{x} \in \mathbb{R}^d$

► Linear predictor :

$$l = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \dots + \theta_d x_d$$

► For convenience of notation, we define  $x_0 = 1$  (then  $\mathbf{x} \in \mathbb{R}^{d+1}$ ) :

$$l = \boldsymbol{\theta}^T \mathbf{x}$$

where  $\boldsymbol{\theta}^T = [\theta_0, \theta_1, \dots, \theta_d] \in \mathbb{R}^{d+1}$  are the parameters of the model

► Binary classification task :  $y \in \{0, 1\}$

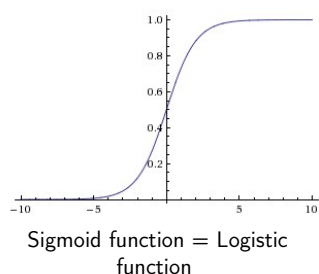
## Logistic regression model

► Hypothesis :  $h_{\boldsymbol{\theta}}(\mathbf{x}) = g(\boldsymbol{\theta}^T \mathbf{x})$

$$\text{with } g(z) = \frac{1}{1 + \exp(-z)}$$

$$h_{\boldsymbol{\theta}}(\mathbf{x}) = \frac{1}{1 + \exp(-\boldsymbol{\theta}^T \mathbf{x})}$$

$$0 \leq h_{\boldsymbol{\theta}}(\mathbf{x}) \leq 1$$



## Interpretation of hypothesis output

$h_{\boldsymbol{\theta}}(\mathbf{x})$  = estimated probability that  $y=1$ , given  $\mathbf{x}$ , parametrized by  $\boldsymbol{\theta}$

$$h_{\boldsymbol{\theta}}(\mathbf{x}) = P(y = 1 | \mathbf{x}, \boldsymbol{\theta}) = \frac{1}{1 + \exp(-\boldsymbol{\theta}^T \mathbf{x})},$$

which implies that :

$$P(y = 0 | \mathbf{x}, \boldsymbol{\theta}) = 1 - P(y = 1 | \mathbf{x}, \boldsymbol{\theta}) = \frac{1}{1 + \exp(\boldsymbol{\theta}^T \mathbf{x})}$$

### Decision boundary

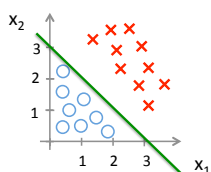
► if  $h_{\boldsymbol{\theta}}(\mathbf{x}) \geq 0.5$ , predict "y=1"

► if  $h_{\boldsymbol{\theta}}(\mathbf{x}) < 0.5$ , predict "y=0"

Because  $g(z) \geq 0.5$  when  $z \geq 0$ ,

$$y = 1 \Leftrightarrow \boldsymbol{\theta}^T \mathbf{x} \geq 0$$

## Decision boundary



$$h_{\boldsymbol{\theta}}(\mathbf{x}) = g(\theta_0 + \theta_1 x_1 + \theta_2 x_2)$$

Predict "y = 1" if  $-3 + x_1 + x_2 \geq 0$

► if  $x_1 + x_2 \geq 3$ , predict "y=1"

► if  $x_1 + x_2 < 3$ , predict "y=0"

How to choose parameters  $\boldsymbol{\theta}$ ?

## Sommaire

Introduction

Principe

Learning by optimization

Going further

## Find parameters $\theta$

### Minimization of a cost function

- We could define :

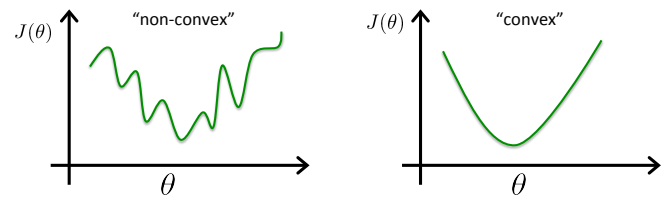
$$J(\theta) = \frac{1}{m} \sum_{i=1}^m \text{loss}(h_{\theta}(\mathbf{x}^{(i)}), y^{(i)})$$

with :

$$\text{loss}(h_{\theta}(\mathbf{x}), y) = \frac{1}{2} (h_{\theta}(\mathbf{x}) - y)^2$$

- But :  $J(\theta)$  non-convex function of the parameters  $\theta$  ! (because  $h_{\theta}(\mathbf{x})$  non-linear function)
- Solution : change the cost function to a convex one.

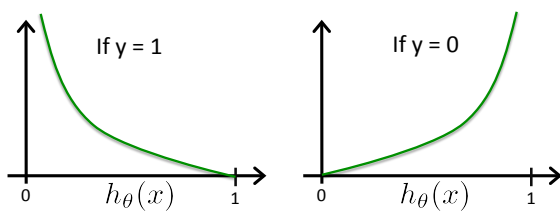
## Convexity



- Convex functions : can be solved quickly and reliably up to very large scale
  - e.g. gradient descent, Lagrange method

## Find parameters $\theta$

$$\text{loss}(h_{\theta}(\mathbf{x}), y) = \begin{cases} -\log(h_{\theta}(\mathbf{x})) & \text{if } y = 1 \\ -\log(1 - h_{\theta}(\mathbf{x})) & \text{if } y = 0 \end{cases}$$



- If  $h_{\theta}(\mathbf{x}) = P(y = 1 | \mathbf{x}, \theta) = 0$ , but  $y = 1$ , the learning algorithm will be penalized with a very large cost.

## Logistic regression cost function

To fit parameters  $\theta$  :  $\min_{\theta} J(\theta)$

$$\begin{aligned} J(\theta) &= \frac{1}{m} \sum_{i=1}^m \text{loss}(h_{\theta}(\mathbf{x}^{(i)}), y^{(i)}) \\ &= -\frac{1}{m} \sum_{i=1}^m y^{(i)} \log(h_{\theta}(\mathbf{x}^{(i)})) + (1 - y^{(i)}) \log(1 - h_{\theta}(\mathbf{x}^{(i)})) \end{aligned}$$

The cost function is :

- convex
- derived from maximum likelihood estimation
- also called **cross-entropy** error function

## Logistic regression cost function

$$\min_{\theta} J(\theta) = -\frac{1}{m} \sum_{i=1}^m y^{(i)} \log(h_{\theta}(\mathbf{x}^{(i)})) + (1 - y^{(i)}) \log(1 - h_{\theta}(\mathbf{x}^{(i)}))$$

### Solution : gradient descent

- $\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta)$
- $\frac{\partial}{\partial \theta_j} J(\theta) = \frac{1}{m} \sum_{i=1}^m (h_{\theta}(\mathbf{x}^{(i)}) - y^{(i)}) x_j^{(i)}$
- simultaneously update all  $\theta_j$

## Sommaire

Introduction

Principle

Learning by optimization

Going further

## Going further

- ▶ Gradient descent parameters and properties
- ▶ Underfitting and adding features to get non linear classifier
- ▶ Overfitting and regularization
- ▶ Multinomial logistic regression

## Problem of underfitting/overfitting

- ▶ **High bias** or **underfitting** is when the form of our hypothesis maps poorly to the trend of the data. It is usually caused by a function that is too simple or uses too few features.
- ▶ At the other extreme, **overfitting** or **high variance** is caused by a hypothesis function that fits the available data but does not generalize well to predict new data. It is usually caused by a complicated function that creates a lot of unnecessary curves and angles unrelated to the data.

## Basis expansion

1

- ▶ Data is likely to be non-linearly separable
- ▶ What if we still wanted to use a linear regression ?
- ▶ How to marry non-linear data to a linear method ?

The trick is to **transform the data** : Map the data onto another features space, such that the data is linear in that space.

→ Including higher order terms **increases the capacity/complexity of the model** : it allows to learn decision boundaries that would be unreachable using simply the original features. This is because a linear decision boundary (which is what logistic regression fits) learned on nonlinear transformations of features will ultimately be nonlinear in terms of the original features.

## Basis expansion

2

- ▶ Denote this transformation  $\varphi : \mathbb{R}^d \rightarrow \mathbb{R}^n$ .
- ▶ If  $\mathbf{x}$  is the original set of features  $\varphi(\mathbf{x})$  denotes the new set of features

Example : Polynomial regression

- ▶ suppose there is just one feature  $x$ .
- ▶ define  $\varphi : \mathbb{R} \rightarrow \mathbb{R}^2$  such that  $\varphi_1(x) = x$  and  $\varphi_2(x) = x^2$
- ▶ the linear predictor becomes  

$$\theta^T \varphi(x) = \theta_0 + \theta_1 \varphi_1(x) + \theta_2 \varphi_2(x) = \theta_0 + \theta_1 x + \theta_2 x^2$$

More generally, a **polynomial basis** is the set of attributes that are powers of  $\mathbf{x}$ .

## Basis expansion

3

- ▶ Data transformation, also known as basis expansion, is a general technique
- ▶ There are many possible choices of  $\varphi$

Example 2 : Radial basis function

A **radial basis function** is a function of the form  $\varphi(\mathbf{x}) = \Phi(\|\mathbf{x} - \mathbf{z}\|)$  where  $\mathbf{z}$  is a constant

- ▶ e.g.  $\varphi(\mathbf{x}) = \|\mathbf{x} - \mathbf{z}\|$  or  $\varphi(\mathbf{x}) = \exp(-\frac{1}{\sigma} \|\mathbf{x} - \mathbf{z}\|^2)$

## Basis expansion

4

- ▶ Basis expansion can significantly increase the utility of methods, especially, linear methods
- ▶ In the above examples, one limitation is that the transformation needs to be defined beforehand
- ▶ One idea is to *learn* the transformation  $\varphi$  from data (e.g., Artificial Neural Networks)
- ▶ Another powerful extension is the use of the *kernel trick* (e.g. SVM)

## Regularization

### Problem of underfitting/overfitting

There are two main options to address the issue of overfitting :

1. Reduce the number of features.
  - ▶ Manually select which features to keep.
  - ▶ Use a model selection algorithm.
2. Regularization
  - ▶ Keep all the features, but reduce the parameters.

## Regularized logistic regression

1

The principle of regularization is to limit the overfitting by simultaneously controlling the model error on the learning set and the values of the model coefficients.

### Intuition

Controlling these coefficients is a way to control the complexity of the model.

This control consists in constraining the coefficients to belong to a subset of  $\mathbb{R}^{d+1}$  rather than being able to take any value in this space. This restricts the set of possible solutions.

Regularization works well when we have a lot of slightly useful features.

## Regularized logistic regression

2

### Cost function

$$J(\theta) = \frac{1}{m} \sum_{i=1}^m \text{loss}(h_{\theta}(\mathbf{x}^{(i)}), y^{(i)}) + \lambda \text{reg}(\theta)$$

- ▶  $\text{loss}(h_{\theta}(\mathbf{x}^{(i)}), y^{(i)}) = y^{(i)} \log(h_{\theta}(\mathbf{x}^{(i)})) + (1 - y^{(i)}) \log(1 - h_{\theta}(\mathbf{x}^{(i)}))$  for logistic regression.
- ▶  $\lambda$  is the **regularization parameter**
  - ▶  $\text{reg}(\theta)$  is a constraint term on the model coefficients  $\theta$
  - ▶  $\lambda$  is an **hyperparameter** of the logistic regression.
  - ▶ If  $\lambda$  is chosen to too large, it may smooth out the function too much and cause underfitting.

## Regularized logistic regression

3

### Ridge regularization (clustering)

$$\text{reg}_r(\theta) = \|\theta\|_2^2 = \sum_{j=1}^d \theta_j^2$$

- ▶  $\sum_{j=1}^d \theta_j^2$  excludes the bias term  $\theta_0$
- ▶ ridge regression uses the  $l_2$  norm of  $\theta$  as a regularizer
- ▶ it has a clustering effect on correlated variables, as correlated variables will have similar coefficients
- ▶ it is a convex optimization problem (quadratic form) that always admits an explicit (analytical) single solution

## Regularized logistic regression

4

### Lasso regularization (sparsity)

$$\text{reg}_l(\theta) = \|\theta\|_1 = \sum_{j=1}^d |\theta_j|$$

- ▶  $\sum_{j=1}^d |\theta_j|$  excludes the bias term  $\theta_0$
- ▶ lasso regression uses the  $l_1$  norm of  $\theta$  as a regularizer
- ▶ it acts as feature selection as it creates a sparse model : some coefficients will be null, leading the corresponding variables to be removed from the model
- ▶ it has no analytical solution, neither always a unique solution, gradient descent should be used.

## Regularized logistic regression

5

### Elastic Net regularization

$$\text{reg}_{el}(\theta) = ((1 - \alpha) \|\theta\|_2^2 + \alpha \|\theta\|_1)$$

- ▶ elastic net regression combines both the  $l_1$  and  $l_2$  norm of  $\theta$  in the regularizer
  - ▶  $l_1$  norm allows to obtain a more easily interpretable model
  - ▶ while  $l_2$  norm avoids the overfitting
- ▶ it is parametrized by  $\alpha \in [0, 1]$

## Higher level view

Logistic regression is a specific type of Generalized Linear Models (GLM).

- ▶ with **binomial** conditional distribution of the response ( $Y$ )
- ▶ parameter  $p = P(Y = 1|X = x)$
- ▶ linear predictor  $\theta^T X$
- ▶ the **logit** function is used to map the linear predictor  $\theta^T X$  to a probability  $p$  :

$$\text{logit}(p) = \log \left[ \frac{p}{1-p} \right] = \theta^T X \Leftrightarrow p = \frac{1}{1 + e^{-\theta^T X}}$$

Because logistic regression predicts *probabilities*, it can be fitted using likelihood.

## Multinomial logistic regression

- ▶ generalizes logistic regression to multiclass problems
- ▶ generalization of the logistic sigmoid : *normalized exponential, softmax function*

$$P(Y = k|\mathbf{x}; \theta) = \frac{\exp(\theta_k^T \mathbf{x})}{\sum_{c=1}^C \exp(\theta_c^T \mathbf{x})}.$$

- ▶ also known as : multiclass LR, softmax regression, multinomial logit, maximum entropy (MaxEnt) classifier, conditional maximum entropy model

## Terminology

### Loss function

$\mathcal{L}(y, h_\theta(\mathbf{x}))$  computes the error for a single training example

- ▶ example : 0/1 loss, hinge loss, cross-entropy loss, exponential loss

### Cost function

usually more general : average of the loss function over the entire training set (empirical risk)

$$C(\theta) = \frac{1}{m} \sum_{i=1}^m \mathcal{L}(y^{(i)}, h_\theta(\mathbf{x}^{(i)}))$$

### Objective function

- ▶ Most general term for any function optimized during the training
- ▶ weighted sum of **cost function** and **regularization**

## Maximum likelihood estimation

Logistic regression assumes a **Bernoulli distribution** defined as :

$$P(Y = 1) = p \text{ and } P(Y = 0) = 1 - p; \text{ with } p \in [0, 1]$$

equivalently :

$$P(Y = y) = p^y (1-p)^{(1-y)}; \text{ for } y \in \{0, 1\}$$

For the logistic regression model :

$$p = h_\theta(\mathbf{x}) = \frac{1}{1 + \exp(-\theta^T \mathbf{x})}$$

Assuming independence of examples in the training set, the likelihood is :

$$L(\theta) = P(y^{(1)}, y^{(2)}, \dots, y^{(m)} | \mathbf{x}^{(1)}, \mathbf{x}^{(2)}, \dots, \mathbf{x}^{(m)}) = \prod_{i=1}^m P(y^{(i)} | \mathbf{x}^{(i)})$$

## Maximum likelihood estimation

$$\begin{aligned} L(\theta) &= \prod_{i=1}^m P(y^{(i)} | \mathbf{x}^{(i)}) \\ &= \prod_{i=1}^m p^{y^{(i)}} (1-p)^{(1-y^{(i)})} \end{aligned}$$

“Log trick” : Instead of maximizing the likelihood, maximise its logarithm (log-likelihood)

$$\begin{aligned} \log L(\theta) &= \log \left( \prod_{i=1}^m P(y^{(i)} | \mathbf{x}^{(i)}) \right) = \sum_{i=1}^m \log(P(y^{(i)} | \mathbf{x}^{(i)})) \\ &= \sum_{i=1}^m y^{(i)} \log(p) + (1 - y^{(i)}) \log(1 - p) \end{aligned}$$

## Cross-entropy

The **negative** log-likelihood for a **single data point** is the **cross-entropy** :

$$-\log(P(y|\mathbf{x})) = -y \log(p) - (1-y) \log(1-p) = H(y, p)$$

Maximise the log-likelihood  $\Leftrightarrow$  Minimise the binary cross-entropy

- ▶ same formula, but two different interpretations

## Exercice - Calcul du gradient

Compute the gradient of the loss function (log-likelihood) for a given example  $\mathbf{x} \in \mathbb{R}^n$  :

$$L(\boldsymbol{\theta}) = L(\theta_0, \theta_1, \dots, \theta_n) = y \log(p) + (1 - y) \log(1 - p)$$

$$\text{with } p = h_{\boldsymbol{\theta}}(\mathbf{x}) = \frac{1}{1 + \exp(-\boldsymbol{\theta}^T \mathbf{x})}$$

and show that :

$$\nabla_{\boldsymbol{\theta}} L(\boldsymbol{\theta}) = [y - h_{\boldsymbol{\theta}}(\mathbf{x})] \mathbf{x}$$