

# Deep Learning

## 4/ Convolutional Neural Network - advanced

---

Francesca Galassi, MCF, Esir

[francesca.galassi@irisa.fr](mailto:francesca.galassi@irisa.fr)

Lab Empenn Irisa-Inria

# ILSVRC : ImageNet Large Scale Visual Recognition Challenge

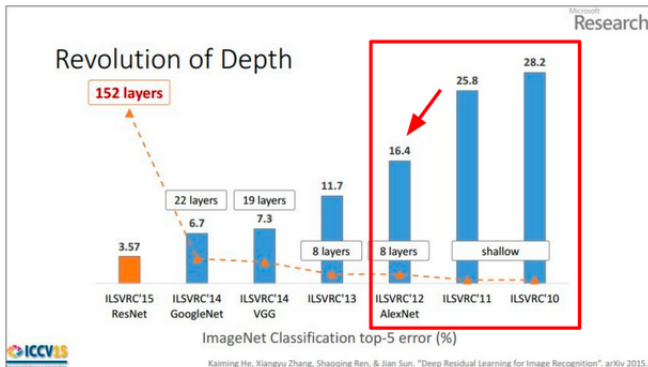
- ⇒ ILSVRC evaluates algorithms for large-scale image classification and object detection since 2010.
- Image classification (1000 categories) and object detection (200 categories), video object detection (30 categories)



*ImageNet dataset* <https://www.image-net.org/>

# ILSVRC : ImageNet Large Scale Visual Recognition Challenge

- ➔ ILSVRC evaluates algorithms for large-scale image classification and object detection since 2010.
  - Image classification (1000 categories) and object detection (200 categories), video object detection (30 categories)
- ➔ ILSVRC 2015 : He et al. introduced **deep residual learning** to train deeper networks effectively.



Slide from He K. presentation at ICCV15


## Residual Block

---


# Residual Block

➡ **Residual blocks** : enabling training of (very) very deep NNs, e.g., over 100 layers

➡ Plain block :



➡ The information flows from  $a^{[\ell]}$  to  $a^{[\ell+2]}$  through all of these steps :




➡ Corresponding equations :

$$\text{layer } [\ell + 1] : \quad z^{[\ell+1]} = W^{[\ell+1]} a^{[\ell]} + b^{[\ell+1]} \quad ; \quad a^{[\ell+1]} = g(z^{[\ell+1]})$$

$$\text{layer } [\ell + 2] : \quad z^{[\ell+2]} = W^{[\ell+2]} a^{[\ell+1]} + b^{[\ell+2]} \quad ; \quad a^{[\ell+2]} = g(z^{[\ell+2]})$$

# Residual Block

➡ **Residual block** :  $a^{[\ell]}$  is injected deeper into the NN via a short cut




$$\text{layer } [\ell + 2] : \quad z^{[\ell+2]} = W^{[\ell+2]} a^{[\ell+1]} + b^{[\ell+2]} \quad ; \quad a^{[\ell+2]} = g(z^{[\ell+2]})$$

➡  $a^{[\ell]}$  is added before applying non-linearity :


$$a^{[\ell+2]} = g(z^{[\ell+2]}) \quad \longrightarrow \quad a^{[\ell+2]} = g(z^{[\ell+2]} + a^{[\ell]})$$

➡ **Residual block** :



# Residual Network

- ⇒ *Plain* very deep network : optimizer has difficulty due to vanishing/exploding gradients
- ⇒ **Residual Network**<sup>1</sup> : stacking residual blocks allows training of very deep networks




- ⇒ ResNets training error keeps decreasing, where *plain* very deep NN tends to increase

1. [He et al., 2015 : Deep Residual networks for image recognition]

# Residual Network

👉 Why do Residual Network work ?



$$\Rightarrow a^{[\ell+2]} = g(z^{[\ell+2]}) \rightarrow a^{[\ell+2]} = g(z^{[\ell+2]} + a^{[\ell]})$$

$$\begin{aligned} a^{[\ell+2]} &= g(z^{[\ell+2]} + a^{[\ell]}) \\ &= g(W^{[\ell+2]} a^{[\ell+1]} + b^{[\ell+2]} + a^{[\ell]}) \end{aligned}$$

$$\Rightarrow \text{If } W^{[\ell+2]} = 0, b^{[\ell+2]} = 0 \implies a^{[\ell+2]} = g(a^{[\ell]}) \xrightarrow{\text{because of ReLU}} a^{[\ell+2]} = a^{[\ell]}$$


Learning the identity function

→ Remark : in ConvNets *same* convolution often used - or add an extra matrix to resize  $a^{[\ell]}$




ResNets<sup>2</sup>

34-layer plain



34-layer residual




2. [He et al., 2015 : Deep Residual networks for image recognition]


## 1x1 Convolution

---

# 1x1 Convolution<sup>3</sup>



→ Simple element-wise multiplication if  $n_c = 1$ .




→ Reducing the number of channels (while introducing non-linearity).


3. [Lin et al. 2013 Network in network]

# Why Use 1x1 Convolution ?

→ Decreasing (increasing) the number of channels  $n_c$  in the input volume :



→ Adding non-linearity, i.e., allowing for more complex function :




## Depthwise Separable Conv

---

# MobileNets

- ⇒ **Low computational cost** at deployment
- ⇒ Ideal for **mobile** and embedded vision applications
- ⇒ Key idea : **depthwise-separable convolutions**<sup>4</sup>
- ⇒ Recall : Standard convolution




- ⇒ Computations :
  - 27 element-wise multiplications ;
  - Add them up ;
  - Shift the filter and repeat.

4. [Howard et al. 2017, MobileNets : Efficient Convolutional Neural Networks for Mobile Vision Applications]

# Standard Convolution vs. Depthwise Separable Convolution

⇒ Standard convolution operation :




⇒ Computational cost = # filter params  $\times$  # filter positions  $\times$  # of filters

=  $3 \times 3 \times 3 \times 4 \times 4 \times 5 = 2160$  computations

# Depthwise Separable Convolution

⇒ **Depthwise separable** convolution :




⇒ Two steps : **Depthwise** Convolution followed by **Pointwise** Convolution



# Depthwise Separable Convolution

## 1. Depthwise convolution :




→ Each filter is applied per channel, resulting in 9 multiplications per filter at 16 locations.

→ The number of output channels is equal to the number of input channels.

$$\begin{aligned}
 \text{Computational cost} &= \# \text{filter params} \times \# \text{filter positions} \times \# \text{of filters} \\
 &= 3 \times 3 \times 4 \times 4 \times 3 = 432
 \end{aligned}$$

# Depthwise Separable Convolution


2. **Pointwise** (1x1) convolution on previous output :




$$\begin{aligned}
 \text{Computational cost} &= \# \text{ filter params} \times \# \text{ filter positions} \times \# \text{ of filters} \\
 &= 1 \times 1 \times 3 \times 4 \times 4 \times 5 = 240 \text{ multiplications}
 \end{aligned}$$

# Depthwise Separable Convolution

⇒ Cost of the standard convolution : 2160 multiplications



⇒ Cost of the Depthwise Separable convolution :  $432 + 240 = 672$  multiplications



# Cost Summary


- ⇒ Cost of standard convolution : 2160
- ⇒ Cost of depthwise separable convolution : 672
- factor of  $\frac{672}{2160} = 0.31 \sim 3$  times more efficient
- ⇒ In a general case<sup>5</sup> :

$$\begin{aligned}
 \frac{\text{depthwise separable convolution}}{\text{standard convolution}} &= \frac{1}{n_c'} + \frac{1}{f^2} \\
 &=_{\text{example}} \frac{1}{5} + \frac{1}{3^2} \\
 &=_{\text{commonly}} \frac{1}{512} + \frac{1}{3^2} \sim 10 \text{ times cheaper}
 \end{aligned}$$

5. [Howard et al. 2017, MobileNets : Efficient Convolutional Neural Networks for Mobile Vision Applications']

# MobileNet


➡ **MobileNet v1** : 13 depthwise separable convolutional layers



➡ Good performance, lower computational cost

# MobileNets

➡ **MobileNet v2**<sup>6</sup> : 17 blocks

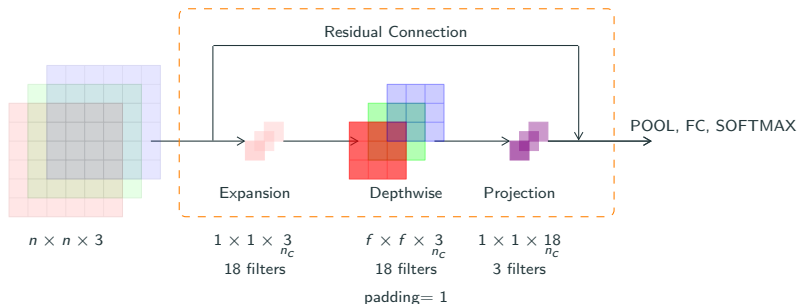


- ➡ Addition of an expansion layer : lower computational cost
- ➡ Addition of a residual (or skip) connection : improved backpropagation of gradients

6. [Sandler et al. 2019, MobileNetV2 : Inverted Residuals and Linear Bottlenecks]

# MobileNets

➡ MobileNet v2 : 17 bottleneck blocks



➡ Given an input  $n \times n \times 3$ , the outputs of the operations in the block are :

- (i.) Expansion output :  $n \times n \times 18$  → expansion by a factor of 6
- (ii.) Depthwise output :  $n \times n \times 18$
- (iii.) Pointwise output :  $n \times n \times 3$  → projection


## Transfer Learning

---



# Transfer Learning : Fine-tuning


- ⇒ Download **open-source** pre-trained weights on a large dataset
- ⇒ **Transfer Learning** : fine-tune the model or adjust the architecture to your data
- ⇒ Example : classification problem with **small training set**
  - ⇒ Download a network for classification, e.g., on ImageNet
  - ⇒ Replace the softmax layer and train only parameters of this layer



# Fine-tuning

Example : classification problem with a **larger training set**.

→ Train the later layers (or replace with your own layers).



Example : classification problem with a **large training set**.

→ Train all network, i.e., use pre-trained weights as initialization.

